



UNIVERSIDAD
DE GRANADA

MUTUAL INFORMATION IN UNSUPERVISED MACHINE LEARNING

FRANCISCO JAVIER SÁEZ MALDONADO

Bachelor's Thesis

Computer Science and Mathematics

Tutor

Nicolás Pérez de la Blanca Capilla

FACULTY OF SCIENCE

H.T.S. OF COMPUTER ENGINEER AND TELECOMMUNICATIONS

Granada, Thursday 21st January, 2021

ABSTRACT

Abstract

CONTENTS

I BASIC CONCEPTS

1	PROBABILITY	5
1.1	Basic notions	5
1.2	Expectation of a random variable	7
2	DISTRIBUTIONS	10

II INFORMATION THEORY

3	MUTUAL INFORMATION	13
3.1	Entropy	13
3.2	Mutual Information	15

III A

Part I

BASIC CONCEPTS

In this part we will introduce the underlying concepts of probability theory and probability distributions that will be needed.

PROBABILITY

Underneath each experiment involving any grade of uncertainty there is a *random variable*. This is no more than a *measurable* function between two *measurable spaces*. A probability space is composed by three elements: $(\Omega, \mathcal{A}, \mathcal{P})$. We will define those concepts one by one.

1.1 BASIC NOTIONS

Definition 1. Let Ω be a non empty sample space. \mathcal{A} is a σ -algebra over Ω if it is a family of subsets of Ω that verify that the emptyset is in \mathcal{A} , and it is closed under complementation and countable unions. That is:

- $\emptyset \in \mathcal{A}$
- If $A \in \mathcal{A}$, then $\Omega \setminus A \in \mathcal{A}$
- If $\{A_i\}_{i \in \mathbb{N}} \in \mathcal{A}$ is a numerable family of \mathcal{A} subsets, then $\cup_{i \in \mathbb{N}} A_i \in \mathcal{A}$

The pair (Ω, \mathcal{A}) is called a *measurable space*. To get to our probability space, we need to define a *measure* on the *measurable space*.

Definition 2. Given (Ω, \mathcal{A}) , a measurable space, a *measure* \mathcal{P} is a countable additive, non-negative set function on this space. That is: $\mathcal{P} : \mathcal{A} \rightarrow \mathbb{R}_0^+$ satisfying:

- $\mathcal{P}(A) \geq \mathcal{P}(\emptyset) = 0$ for all $A \in \mathcal{A}$
- $\mathcal{P}(\cup_n A_n) = \sum_n \mathcal{P}(A_n)$ for any countable collection of disjoint sets $A_n \in \mathcal{A}$.

If $\mathcal{P}(\Omega) = 1$, \mathcal{P} is a *probability measure* or simply a *probability*. With the concepts that have just been explained, we get to the following definition:

Definition 3. A *measure space* is the tuple $(\Omega, \mathcal{A}, \mathcal{P})$ where \mathcal{P} is a *measure* on (Ω, \mathcal{A}) . If \mathcal{P} is a *probability measure* $(\Omega, \mathcal{A}, \mathcal{P})$ will be called a *probability space*.

Throughout this work, we will be always in the case where \mathcal{P} is a probability measure, so we will always be talking about probability spaces. Some notation for these measures must be introduced. Let A and B be two events. The notation $P(A, B)$ refers to the probability of the intersection of the events A and B , that is: $P(A, B) := P(A \cap B)$. It is clear that since $A \cap B = B \cap A$, then $P(A, B) = P(B, A)$. We remark the next definition since it will be important.

Definition 4. Let A, B be two events in Ω . The *conditional probability* of B given A is defined as:

$$P(B|A) = \frac{P(A, B)}{P(A)}$$

There is an alternative way to state the definition that we have just made.

Theorem 1 (Bayes' theorem). Let A, B be two events in Ω , given that $P(B) \neq 0$. Then

$$P(B|A) = \frac{P(A|B)P(B)}{P(A)}$$

Proof. Straight from the definition of the conditional probability we obtain that:

$$P(A, B) = P(A|B)P(B)$$

We also see from the definition that

$$P(B, A) = P(B|A)P(A)$$

Hence, since $P(A, B) = P(B, A)$,

$$P(A|B)P(B) = P(B|A)P(A) \implies P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

□

However, events might not give any information about another event occurring. When this happens, we call those events to be *independent*. Mathematically, if A, B are independent events:

$$P(A, B) = P(A)P(B)$$

and as a consequence of this, the conditional probability of those events is $P(A|B) = P(A)$. For a finite set of events $\{A_i\}_{i=1}^n$, we say that they are mutually independent if and only if every event is independent of any intersection of the other events. That is, if $\{B_i\} \subset \{A_i\}$, then

$$P\left(\bigcap_{i=1}^k B_i\right) = \prod_{i=1}^k P(B_i) \quad \forall k \leq n$$

Random variables (R.V.) can now be introduced. Their first property is that they are measurable functions. These kind of functions are defined as it follows:

Definition 5. Let $(\Omega_1, \mathcal{A}), (\Omega_2, \mathcal{B})$ be measurable spaces. A function $f : \Omega_1 \rightarrow \Omega_2$ is said to be *measurable* if, $f^{-1}(B) \in \mathcal{A}$ for every $B \in \mathcal{B}$.

As a quick note, we can affirm that if f, g are real-valued measurable functions, and $k \in \mathbb{R}$, it is true that $kf, f + g, fg$ and f/g (if g is not the identically zero function) are also *measurable functions*.

We are now ready to define one of the concepts that will lead us to the main objective of this thesis.

Definition 6 (Random variable). Let $(\Omega, \mathcal{A}, \mathcal{P})$ be a probability space, and (E, \mathcal{B}) be a measurable space. A *random variable* is a measurable function $X : \Omega \rightarrow E$, from the probability space to the measurable space. This means: for every subset $B \in (E, \mathcal{B})$, its preimage

$$X^{-1}(B) = \{\omega : X(\omega) \in B\} \in \mathcal{A}.$$

Using that sums, products and quotients of measurable functions are measurable functions, we obtain that *sums, products and quotients of random variables are random variables*.

Let now X be a R.V. The *probability* of X taking a concrete value on a measurable set contained in E , say, $S \in E$, is written as:

$$P_X(S) = P(X \in S) = P(\{a \in \Omega : X(a) \in S\})$$

A very simple example of random variable is the following:

Example 1. Consider tossing a coin. The possible outcomes of this experiment are *Heads* or *Tails*. Those are our random events. We can give our random events a possible value. For instance, let *Heads* be 1 and *Tails* be 0. Then, our random variable looks like this:

$$X = \begin{cases} 1 & \text{if we obtain heads} \\ 0 & \text{if we obtain tails} \end{cases}$$

In the last example, our random variable is *discrete*, since the set $\{X(\omega) : \omega \in \Omega\}$ is finite. A *Random Variable* can also be *continuous*, if it can take any value within an interval.

1.2 EXPECTATION OF A RANDOM VARIABLE

Definition 7. The *cumulative distribution function* F_X of a real-valued random variable X is its probability of taking value below or equal to x . That is:

$$F_X(x) = P(X \leq x) = P(\{\omega : X(\omega) \leq x\}) = P_X((-\infty, x]) \quad \forall x \in \mathbb{R}$$

Depending on the image of a random variable X , we can difference between certain types of random variables. If the image \mathcal{X} of X is countable, we call it a *discrete* random variable. Its *probability mass function* gives the probability of the r.v. being equal to a certain value:

$$p(x) = P(X = x).$$

If the image \mathcal{X} of X is uncountable and real, then X is a *continuous* random variable. In this case there might exist a non-negative Lebesgue-integrable function f such that:

$$F_X(x) = \int_{-\infty}^x f(t)dt,$$

called the *probability density function* of X .

We are now ready to introduce the *expectation* of a random variable.

Definition 8 (Expectation of a R.V.). Let X be a non negative random variable on a probability space $(\Omega, \mathcal{A}, \mathcal{P})$. The expectation $E[X]$ of X is defined as:

$$E[X] = \int_{\Omega} X(\omega) dP(\omega)$$

The expectation of a random variable will be also denoted as μ . Now, if X is generic R.V, the expectation is defined as:

$$E[X] = E[X^+] - E[X^-]$$

where X^+, X^- are defined as it follows:

$$X^+(\omega) = \max(X(\omega), 0) \quad X^-(\omega) = \min(X(\omega), 0)$$

The *expectation* $E[X]$ of a *random variable* is a linear operation. That is, if \mathcal{Y} is another random variable, and $\alpha, \beta \in \mathbb{R}$, then

$$E[\alpha X + \beta \mathcal{Y}] = \alpha E[X] + \beta E[\mathcal{Y}]$$

this is a trivial consequence of the linearity of the *Lebesgue integral*.

As a note, if X is a *discrete* random variable and \mathcal{X} is its image, its expectation can be computed as:

$$E[X] = \sum_{x \in \mathcal{X}} x P_X(x)$$

where x is each possible outcome of the experiment, and $P_X(x)$ the probability under the distribution of X of the outcome x . The expression given in the definition before generalizes this particular case.

Using the definition of the *expectation* of a random variable, we can approach to the *moments* of a random variable.

Definition 9. If $k \in \mathbb{N}$, then $E[X^k]$ is called the k – *th* moment of X .

If we take $k = 1$, we have the definition of the *expectation*. It is sometimes written as $m_X = E[X]$, and called the *mean*. We use the *mean* in the definition of the variance:

Definition 10. Let X be a random variable. If $E[X^2] < \infty$, then the *variance* of X is defined to be

$$\text{Var}(X) = E[(X - m_X)^2] = E[X^2] - m_X^2$$

Thanks to the linearity of the *expectation* of a random variable, it is easy to see that, if $a, b \in \mathbb{R}$, then

$$\text{Var}(aX + b) = E[(aX + b) - E[aX + b]]^2 = a^2 E[(X - m_X)^2] = a^2 \text{Var}(X)$$

Usually, when it comes to applying these concepts to a real problem, we will be looking at multiple variables. We would like to have a collection of random variables each one representing one of this variables. In order to set the notation for these kinds of situations, we will introduce *random vectors*.

Definition 11. A random vector is a row vector $\mathbf{X} = (X_1, \dots, X_n)$ whose components are real-valued random variables on the same probability space (Ω, \mathcal{A}, P) .

The probability distribution of a random variable can be extended in to the *joint probability distribution* of a random vector.

Definition 12. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random vector. The *cumulative distribution function* $F_{\mathbf{X}} : \mathbb{R}^n \rightarrow [0, 1]$ of \mathbf{X} is defined as:

$$F_{\mathbf{X}}(x) = P(X_1 \leq x_1, \dots, X_n \leq x_n).$$

We also name it *multivariate distribution*. We explained before the independence of a pair of events. Using the cumulative distribution function, we can now define the independence between random variables.

Definition 13. A finite set of n random variables $\{X_1, \dots, X_n\}$ is mutually independent if and only if, for any sequence $\{x_1, \dots, x_n\}$, the events $\{X_1 \leq x_1\}, \dots, \{X_n \leq x_n\}$ are mutually independent. Equivalently, this finite set is mutually independent if and only if:

$$F_{X_1, \dots, X_n}(x_1, \dots, x_n) = F_{X_1}(x_1) \dots F_{X_n}(x_n) \quad \text{for all } x_1, \dots, x_n.$$

We can also extend the notion of expectation to a random vector. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random vector and assume that $E[X_i]$ exists for all $i \in \{1, \dots, n\}$. The expectation of \mathbf{X} is defined as the vector containing the expectations of each individual random vector, that is:

$$E[\mathbf{X}] = \begin{bmatrix} E[X_1] \\ \vdots \\ E[X_n] \end{bmatrix}$$

To generalize the variance of a random variable, we have to build the following matrix.

Definition 14. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random vector. Then, the *covariance matrix* of \mathbf{X} is defined as:

$$\Sigma = \text{Cov}(\mathbf{X}) = E[(\mathbf{X} - \mu_{\mathbf{X}})(\mathbf{X} - \mu_{\mathbf{X}})^T] = \begin{pmatrix} \sigma_{11} & \cdots & \sigma_{1n} \\ \vdots & \ddots & \vdots \\ \sigma_{n1} & \cdots & \sigma_{nn} \end{pmatrix}$$

where $\sigma_{ij} = \text{Cov}(X_i, X_j) = E[(X_i - \mu_i)(X_j - \mu_j)] = \sigma_{ji}$.

It can also happen that, given a *random vector*, we would like to know the probability distribution of a few of its components. That is called the *marginal distribution*.

Definition 15 (Marginal Distribution). Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random vector. The marginal distribution of a subset of \mathbf{X} is the probability distribution of the variables contained in the subset.

In the simple case of having two random variables, e.g. X_1 and X_2 , then the marginal distribution of X_1 is:

$$P(x) = \int_{X_2} P(x_1, x_2) dx_2.$$

DISTRIBUTIONS

We have introduced the concepts of *random variable*, *random vector* and its *probability distribution*. Now, given two distributions, in the following chapters we will like to see how different they are from each other. In order to compare them, we enunciate the definition of the Kullback-Leibler divergence.

Definition 16 (Kullback-Leibler Divergence). Let P and Q be probability distributions over the same probability space Ω . Then, the Kullback-Leibler divergence is defined as:

$$D_{KL}(P \parallel Q) = E_P \left[\log \frac{P(x)}{Q(x)} \right]$$

It is defined if and only if P is *absolutely continuous with respect to* Q , that is, if $P(A) = 0$ for any A subset of Ω where $Q(A) = 0$. There are some properties of this definition that must be stated. The first one is the following proposition:

Proposition 1. If P, Q are two probability distributions over the same probability space, then $D_{KL}(P \parallel Q) \geq 0$.

Proof. Firstly, note that if $a \in \mathbb{R}^+$, then $\log a \leq a - 1$. Then:

$$\begin{aligned} -D_{KL}(P \parallel Q) &= -E_P \left[\log \frac{P(x)}{Q(x)} \right] \\ &= E_P \left[\log \frac{Q(x)}{P(x)} \right] \\ &\leq E_P \left[\left(\frac{Q(x)}{P(x)} - 1 \right) \right] \\ &= \int P(x) \frac{Q(x)}{P(x)} dx - 1 \\ &= 0 \end{aligned}$$

So we have obtained that $-D_{KL}(P \parallel Q) \leq 0$, which implies that $D_{KL}(P \parallel Q) \geq 0$. \square

As a corollary of this proposition, we can affirm that $D_{KL}(P \parallel Q)$ equals zero if and only if $P = Q$ almost everywhere. We will also remark the discrete case, as it will be used later. Let P, Q be discrete probability distributions defined on the same probability space Ω . Then,

$$D_{KL}(P \parallel Q) = \sum_{x \in \Omega} P(x) \log \left(\frac{P(x)}{Q(x)} \right)$$

Part II

INFORMATION THEORY

Information theory is the base for all the following work. In this part, *Mutual Information* will be explained and then, bounds for this function will be given.

MUTUAL INFORMATION

Obtaining good representations of data is one of the most important tasks in Machine learning. Recently, it has been discovered that maximizing *Mutual Information* between two elements in our data can give us good representations for our data. We will go through the basic concepts first.

3.1 ENTROPY

The *mutual information* concept is based on the *Shannon entropy*, which we will introduce first, along with some basic properties of it. The *Shannon entropy* is a way of measuring the uncertainty in a random variable. Given an event $\mathcal{A} \in \Omega$, P a probability measure and $P[\mathcal{A}]$ the probability of \mathcal{A} , we can affirm that

$$\log \frac{1}{P[\mathcal{A}]}$$

describes *how surprising is that \mathcal{A} occurs*. For instance, if $P[\mathcal{A}] = 1$, then the last expression is zero, which means that it is not a surprise that \mathcal{A} occurred. With this motivation, we get to the following definition.

Definition 17. Let X be a discrete random variable with image \mathcal{X} . The *Shannon entropy*, or simply *entropy*, $H(X)$ of X is defined as:

$$H(X) = E_X \left[\log \frac{1}{P_X(X)} \right] = \sum_{x \in \mathcal{X}} P_X(x) \log \frac{1}{P_X(x)}$$

The *entropy* can trivially be expressed as:

$$H(X) = - \sum_{x \in \mathcal{X}} P_X(x) \log P_X(x)$$

There are some properties of the *entropy* that must be remarked.

Proposition 2. Let X be a random variable with image \mathcal{X} . then

$$0 \leq H(X) \leq \log(|\mathcal{X}|)$$

Proof. Since $\log y$ is concave on \mathbb{R}^+ , by Jensen's inequality, see ?? 4.:

$$H(X) = - \sum_{x \in X} P_X(x) \log P_X(x) \leq \log \left(\sum_{x \in \mathcal{X}} 1 \right) = \log(|\mathcal{X}|)$$

For the lower bound, it is easy to see that, since $P_X(x) \in [0, 1] \ \forall x \in \mathcal{X}$, $\log P_X(x) \leq 0 \ \forall x \in \mathcal{X}$. Then, $-P_X(x) \log P_X(x) \geq 0$ for all $x \in X$, so $H(X) \geq 0$. \square

We can also see that the equality on the left holds if and only if exists x in X such that its probability is exactly one, that is $P_X(x) = 1$. The right equality holds if and only if, for all $x \in \mathcal{X}$, its probability is $P_X(x) = \frac{1}{|\mathcal{X}|}$.

Conditional entropy

We have already said that *entropy measures* how surprising is that an event occurs. Usually, we will be looking at two random variables and it would be interesting to see how surprising is that one of them, say X , occurred, if we already know that Y occurred. This leads us to the definition of *conditional entropy*. Lets see a simpler case first:

Let A be an event, and X a random variable. The conditional probability $P_{X|A}$ defines the entropy of X conditioned to A :

$$H(X|A) = \sum_{x \in \mathcal{X}} P_{X|A}(x) \log \frac{1}{P_{X|A}(x)}$$

If Y is another random variable and \mathcal{Y} is its image, intuitively we can sum the conditional entropy of an event with all the events in \mathcal{Y} , and this way we obtain the conditional entropy of X given Y .

Definition 18 (Conditional Entropy). Let X, Y be random variables with images \mathcal{X}, \mathcal{Y} . The *conditional entropy* $H(X|Y)$ is defined as:

$$\begin{aligned} H(X|Y) &:= \sum_{y \in \mathcal{Y}} P_Y(y) H(X|Y = y) \\ &= \sum_{y \in \mathcal{Y}} P_Y(y) \sum_{x \in \mathcal{X}} P_{X|Y}(x|y) \log \frac{1}{P_{X|Y}(x|y)} \\ &= \sum_{x \in \mathcal{X}, y \in \mathcal{Y}} P_{XY}(x, y) \log \frac{P_Y(y)}{P_{XY}(x, y)} \end{aligned}$$

The interpretation of the *Conditional Entropy* is simple: the uncertainty in X when Y is given. Since we know about an event that has occurred (Y), intuitively the conditional entropy, or the uncertainty of X occurring given that Y has occurred, will be lesser than the entropy of X , since we already have some information about what is happening. We can prove this:

Proposition 3. Let X, Y be random variables with images \mathcal{X}, \mathcal{Y} . Then:

$$0 \leq H(X|Y) \leq H(X)$$

Proof. The inequality on the left was proved on Proposition ?? 2. The characterization of when $H(X|Y) = 0$ was also mentioned after it. Let's look at the inequality on the right. Note that, restricting to the (x, y) where $P_{XY}(x, y) > 0$ and using the definition of the conditional probability we have:

$$\begin{aligned} H(X|Y) &= \sum_y P_Y(y) \sum_x P_{X|Y}(x|y) \log \frac{1}{P_{X|Y}(x|y)} \\ &= \sum_{x,y} P_Y(y) P_{X|Y}(x|y) \log \frac{P_Y(y)}{P_{XY}(x, y)} = \sum_{x,y} P_{XY}(x, y) \log \frac{P_Y(y)}{P_{XY}(x, y)} \end{aligned}$$

and

$$H(X) = \sum_x P_X(x) \log \frac{1}{P_X(x)} = \sum_{x,y} P_{XY}(x, y) \log \frac{1}{P_X(x)}$$

hence,

$$H(X|Y) - H(X) = \sum_{x,y} P_{XY}(x, y) \left(\log \frac{P_Y(y)}{P_{XY}(x, y)} - \log \frac{1}{P_X(x)} \right) = \sum_{x,y} P_{XY} \log \frac{P_Y(y) P_X(x)}{P_{XY}(x, y)} \quad (1)$$

so, using Jensen's Inequality , we obtain:

$$\begin{aligned} \sum_{x,y} P_{XY} \log \frac{P_Y(y) P_X(x)}{P_{XY}(x, y)} &\leq \log \left(\sum_{x,y} \frac{P_{XY}(x, y) P_Y(y) P_X(x)}{P_{XY}(x, y)} \right) \\ &= \log \left(\left(\sum_x P_X(x) \right) \left(\sum_y P_Y(y) \right) \right) = \log 1 = 0 \end{aligned}$$

and this leads us to:

$$H(X|Y) - H(X) \leq 0 \implies H(X|Y) \leq H(X)$$

as we wanted. \square

It must be noted that, on the development of $H(X|Y) - H(X)$, in the first inequality, equality holds if and only if $P_{XY}(x, y) = P_X(x)P_Y(y)$ for all (x, y) with $P_{XY}(x, y) > 0$, as it is said in Jensen's inequality. For the second inequality, equality holds if and only if $P_{XY}(x, y) = 0$, which implies $P_X(x)P_Y(y) = 0$ for any $x \in \mathcal{X}, y \in \mathcal{Y}$. It follows that $H(X|Y) = H(X)$ if and only if $P_{XY}(x, y) = P_X(x)P_Y(y)$ for all $(x, y) \in \mathcal{X} \times \mathcal{Y}$.

3.2 MUTUAL INFORMATION

Using the *entropy* of a random variable we can directly state the definition of *Mutual Information* as it follows:

Definition 19 (Mutual Information). Let X, Z be random variables. The *Mutual Information* (MI) is expressed as the difference between the entropy of X and the conditional entropy of X and Z , that is:

$$I(X, Z) := H(X) - H(X|Z)$$

Since the entropy of the random variable $H(X)$ explains the uncertainty of X occurring, the intuitive idea of the MI is to determine the decrease of uncertainty of X occurring when we already know that Z has occurred. We also have to note that, using the definition of the *entropy* and the expression obtained in 1, we can rewrite the MI as it follows:

$$\begin{aligned} I(X, Z) &= \sum_{x \in \mathcal{X}} P_X(x) \log \frac{1}{P(x)} - \sum_{x \in \mathcal{X}, z \in \mathcal{Z}} P_{XZ}(x, z) \log \frac{P_Z(z)}{P_{XZ}(x, z)} \\ &= \sum_{x, z} P_{XZ} \log \frac{P_Z(z) P_X(x)}{P_{XZ}(x, z)} = D_{KL}(P_{XZ} \parallel P_X P_Z) \end{aligned}$$

and we have obtained an expression of the mutual information using the *Kullback-Leibler* divergence. This provides with the following immediate consequences:

- (i) Mutual information is non-negative, that is: $I(X, Z) \geq 0$.
- (ii) If X, Z are random variables, then its mutual information equals zero if and only if they are independent. This is trivial because if $D_{KL}(P_{XZ} \parallel P_X P_Z) = 0$, then $P_{XZ} = P_X P_Z$ almost everywhere so X and Z are independent.
- (iii) Since $P_{XZ} = P_{ZX}$ and $P_X P_Z = P_Z P_X$, mutual information is symmetric. That is: $I(X, Z) = I(Z, X)$.

Part III

A

Appendices

This appendix will be used to set forth some theoretical results that might not always be relevant but are needed to understand some details during this thesis. Not all of them will be proven.

Proposition 4 (Jensen's Inequality). *Let $f : \mathcal{D} \rightarrow \mathbb{R}$ be a concave function and $n \in \mathbb{N}$. For any $p_1, \dots, p_n \in \mathbb{R}_0^+$ with $\sum p_i = 1$ and any $x_1, \dots, x_n \in \mathbb{D}$, it holds that:*

$$\sum_{i=1}^n p_i f(x_i) \leq f\left(\sum_{i=1}^n p_i x_i\right)$$

Furthermore, if f is strictly concave and $p_i \geq 0$ for all $i = 1, \dots, n$, then the equality holds if and only if $x_1 = \dots = x_n$

BIBLIOGRAPHY

Löwe, Sindy, O'Connor, Peter, & Veeling, Bastiaan. 2019. Putting an End to End-to-End: Gradient-Isolated Learning of Representations. *Pages 3039–3051 of: Advances in Neural Information Processing Systems*.